

# Solving the Curse of Dimensionality in DDCM

Graph-Based Computation for Activity-Based Travel Demand Models

January 2026

Press Space for next page →

# 1. Why This Matters

The Problem We're Solving

# The Policy Questions That Motivate Us



Urban mobility: trams, cars, buses, pedestrians competing for space

## 🚌 Transit Planning

"If we add a new bus line, who will use it?"

## 🏡 Remote Work

"How will work-from-home change commute patterns?"

## 💰 Congestion Pricing

"Who benefits and who loses from road pricing?"

**Challenge:** Model how people make daily activity and travel decisions.

# What Are Travel Demand Models?

## ANSWERING POLICY QUESTIONS

To answer questions like "Who will use a new bus line?" or "How will remote work change commutes?", we need models that **predict how people make travel decisions**.

### What They Do

-  Predict **where** people travel
-  Predict **when** they travel
-  Predict **how** they travel (mode choice)
-  Forecast impacts of policy changes

### The Key Insight

Travel is a **derived demand**—people don't travel for fun, they travel to participate in *activities* (work, shopping, leisure).

→ Modeling **activity choice** gives us the complete picture.

# What is DDCM?

Dynamic Discrete Choice Model (Västberg et al., 2020 — *Transportation Science*)

## The Core Idea

Your daily schedule is a **sequence of decisions**:

- When should I leave home?
- Should I go to work or run errands first?
- How long should I stay at the office?
- What's the best way to get there?

## Key Properties

- ✓ **Economic foundation** — enables welfare analysis
- ✓ **Time consistency** — respects physical constraints
- ✓ **Interdependence** — choices affect future options

**Case Study:** Workdays in Stockholm, Sweden — 1,240 zones × 4 modes × 6 activities

# DDCM as a Markov Decision Process (MDP)

An MDP  $(S, C, q, u)$  models sequential decision-making under uncertainty.

**State ( $S$ ) — "Where am I now?"**

Your current situation:

Time, location, activity, mode, history

**Action ( $C$ ) — "What can I do?"**

Available choices:

Stay, travel to destination, change activity

**Transition ( $q$ ) — "What happens next?"**

How actions move you to new states:

Current state + action → next state

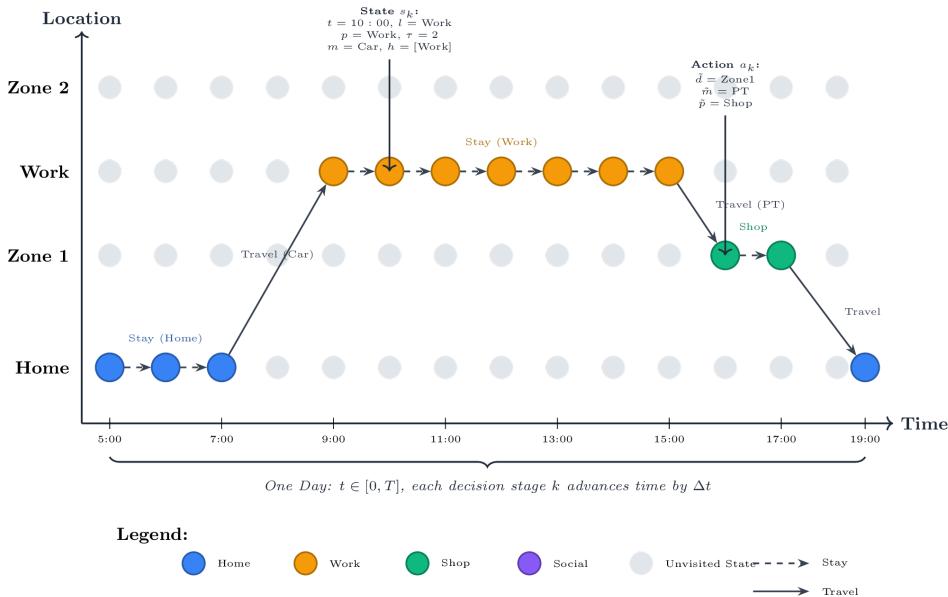
**Utility ( $u$ ) — "How good is this?"**

The "reward" for each choice:

Activity enjoyment – travel discomfort

**Goal:** Find policy  $\pi(s)$  that maximizes *total expected utility* over the day.

# Visualizing a Daily Activity Pattern

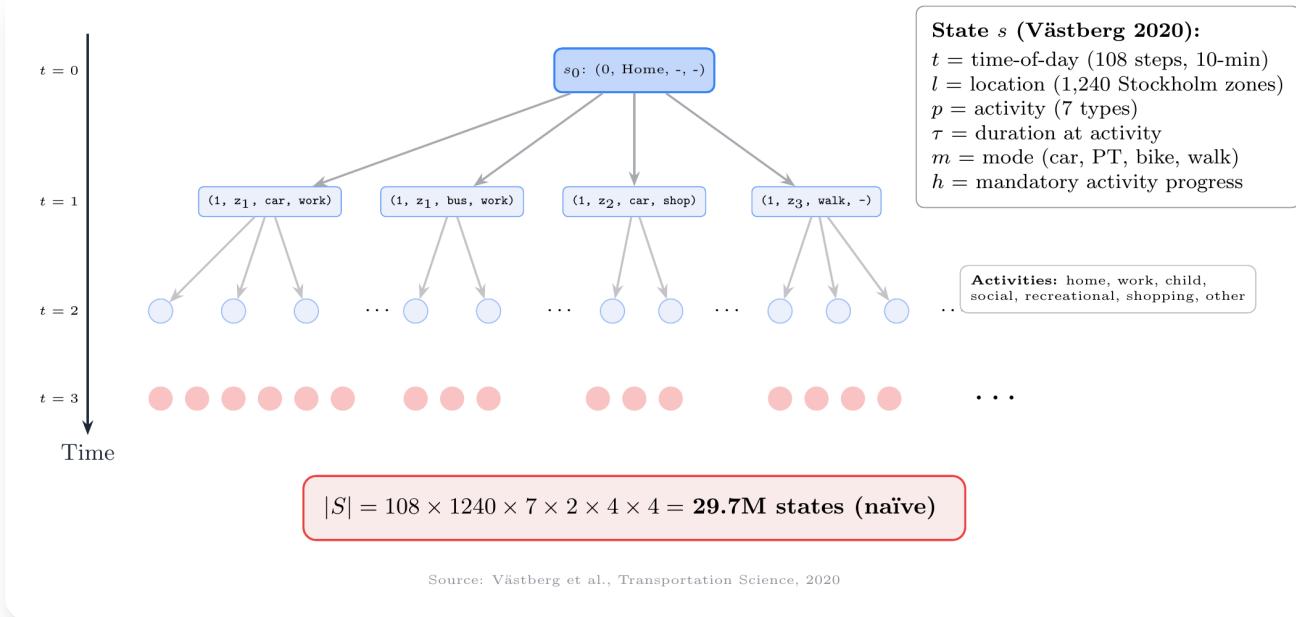


A typical workday: Morning at Home → Travel to Work → Work → Travel to Shopping → Return Home

# 2. The Problem

The Curse of Dimensionality

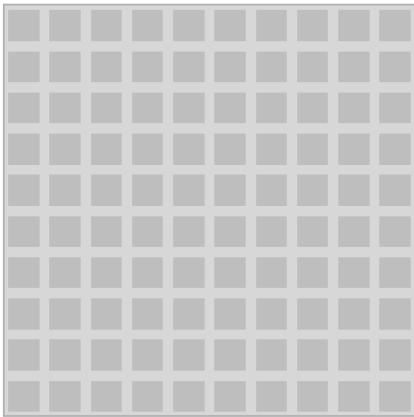
# The Curse of Dimensionality



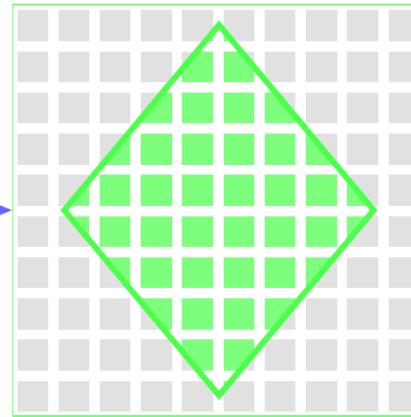
Exponential growth of state space over time

# The Key Insight

Theoretical State Space



Reachable States



Prune

*“What if we don’t compute the impossible ones?”*

# Research Objectives

## PRIMARY OBJECTIVE

Make DDCM **computationally tractable** for population-scale simulation.

## TARGET OUTCOMES

- ✓ Reduce computational burden by orders of magnitude
- ✓ Enable simulation on consumer hardware (GPU)
- ✓ Preserve theoretical properties

## THE CHALLENGE

Current: ~10 seconds per agent → 100,000 agents = **1,000+ CPU-days**

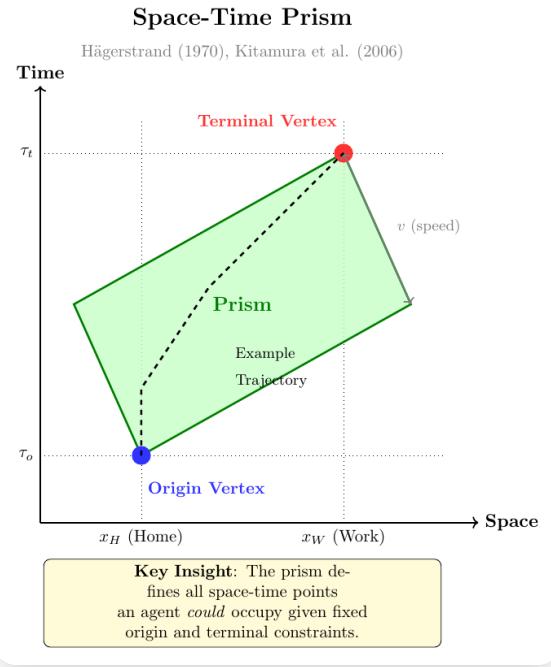
Goal: Make this tractable while maintaining **optimality**

# 3. Theoretical Foundation

The Reachability Concept

# The Space-Time Prism

Hägerstrand (1970): Not everywhere is reachable from everywhere



## KEY CONCEPT

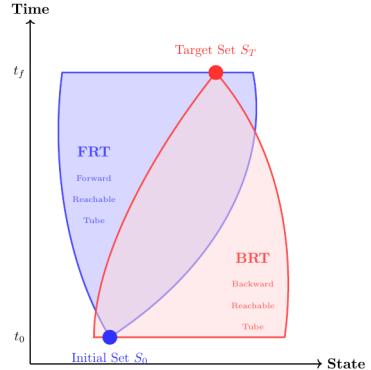
- **Origin:** When/where you can leave
- **Destination:** When/where you must arrive
- **Speed Constraint:** Limits how far you can travel
- **The Prism:** All space-time points you *could* occupy

Activity scheduling has inherent constraints!

# Reachability Analysis

## Hamilton-Jacobi Reachability

Control Theory / Robotics

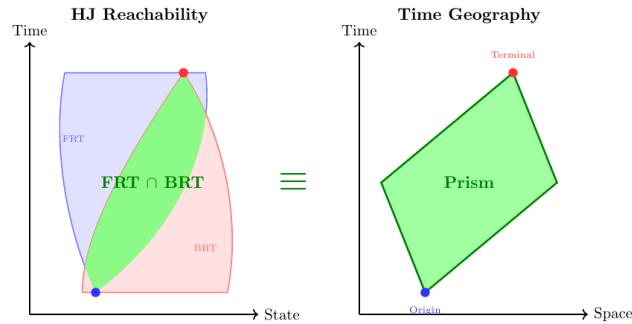


FRT: All states reachable FROM initial set  
BRT: All states that CAN REACH target set

Forward and Backward Reachable Tubes

## Connecting Concepts

$\text{Prism} \equiv \text{FRT} \cap \text{BRT}$



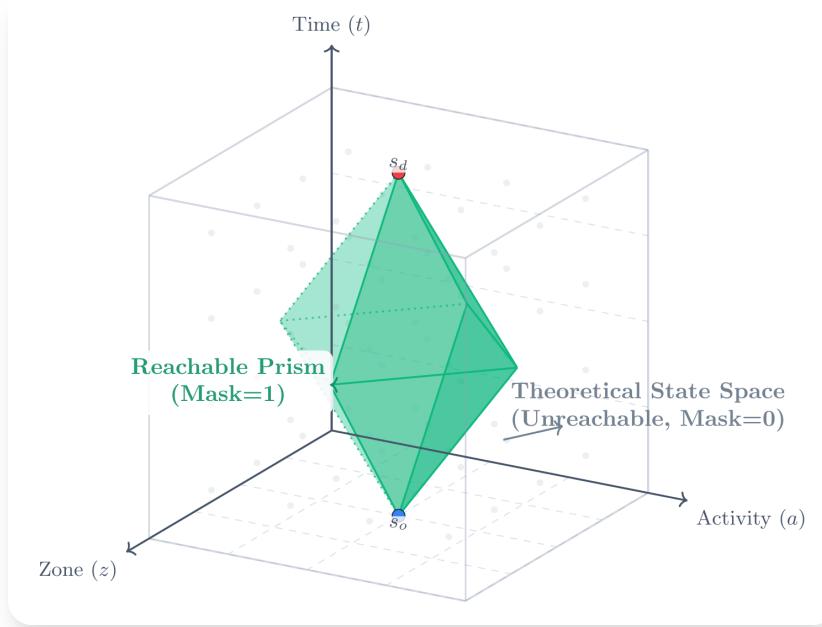
The Insight: Both concepts describe the same thing—the feasible region where trajectories can start AND reach destination.

Our DDCM builds this region explicitly as a sparse graph.

Intersection = Feasible Region

$\text{Prism} \equiv \text{FRT} \cap \text{BRT}$  — Only compute where you CAN go AND CAN return from!

# Only Compute the Reachable



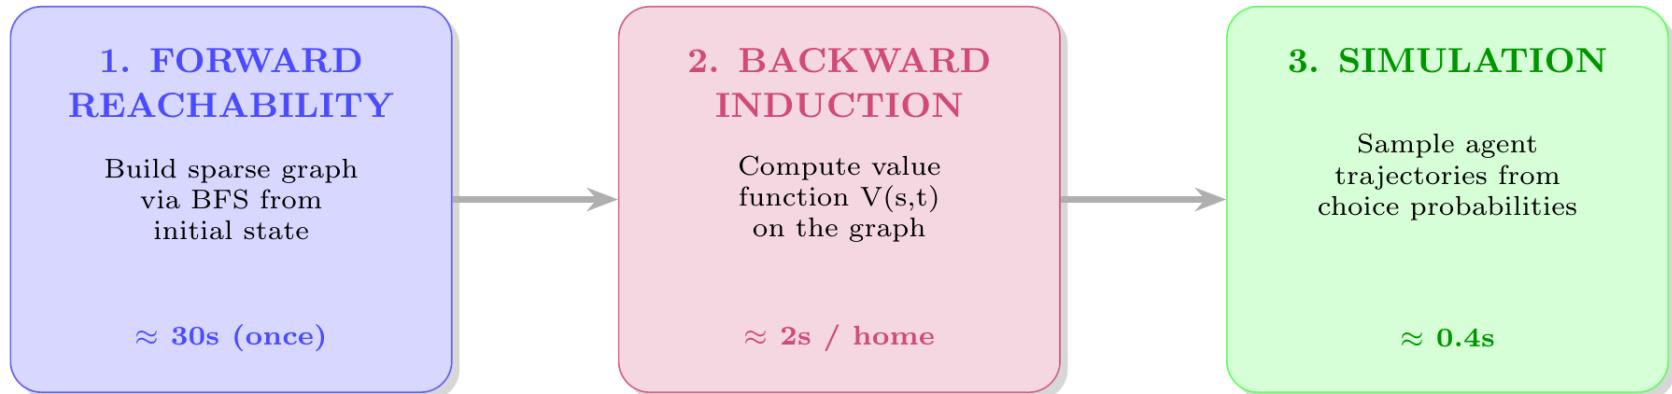
Green Prism (Reachable States) inside Gray Cube (Theoretical State Space)

# 4. Our Solution

The Three-Phase Pipeline

# The Three-Phase Pipeline

## The Three-Phase Pipeline



*GPU-accelerated — Reachability-pruned — Optimal solutions*

# How It Works: Value Computation

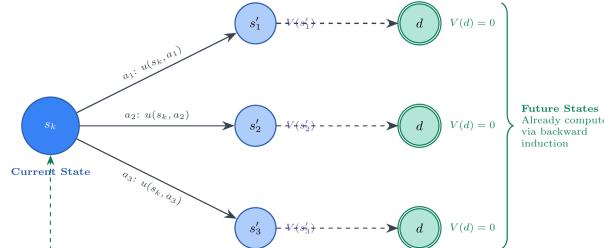
## THE CORE EQUATION

At each state, the agent chooses the action that maximizes:

Immediate utility + Future utility

This is the Bellman equation from dynamic programming.

### The Bellman Equation: Value Function Propagation



#### Expected Value Function at $x_k$

$$\hat{V}(x_k) = \log \sum_{a_k \in C(x_k)} e^{u(x_k, a_k) + EV(x_k, a_k)}$$

#### Choice Probability

$$P(a_k|x_k) = \frac{e^{u(x_k, a_k) + EV(x_k, a_k)}}{\sum_{\tilde{a}_k} e^{u(x_k, \tilde{a}_k) + EV(x_k, \tilde{a}_k)}}$$

#### Key Insight

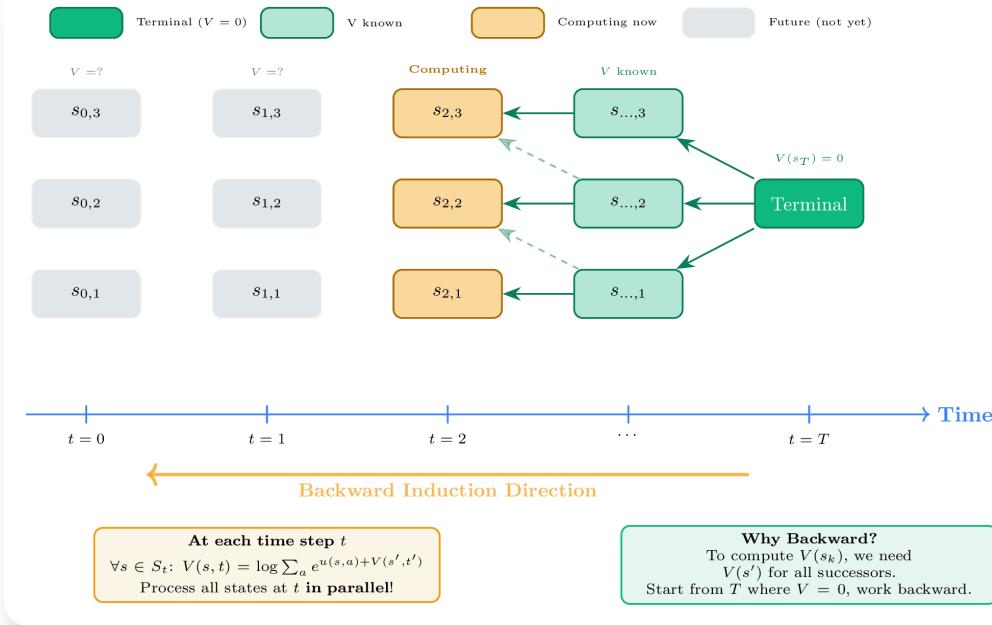
Each action's total value =  
Immediate utility  $u(x_k, a_k)$   
+ Expected future  $EV(x_k, a_k)$

#### LogSumExp = Soft Max

Not just "pick best action"  
but account for *option value*  
of all possibilities

**Key property:** The "log-sum" formula gives both choice probabilities AND consumer surplus.

# Working Backward in Time



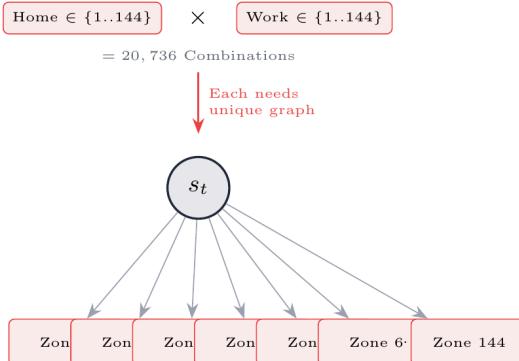
Start at end of day (known value)  $\rightarrow$  Work backward  $\rightarrow$  Future values always available when needed

# The Heterogeneity Challenge

Different agents have different home and work locations. Each (Home, Work) pair has a **unique graph structure**.

## Concrete Approach

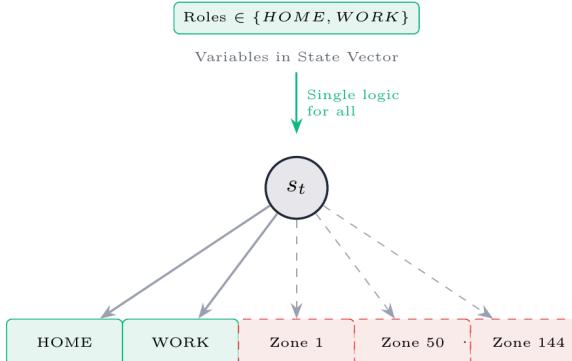
Input: Agent Heterogeneity



**20,736 Unique Trees**  
Structure depends on specific  $(H, W)$ .  
Total: **316+ HOURS**

## RMDP Approach

Input: Abstract Roles



**1 Universal Tree**  
Structure is identical for everyone.  
Total: **47 SECONDS**

\*Roles resolved at runtime via ZoneBinding  
(e.g., HOME  $\rightarrow$  Zone 50)

# RMDP: The Key Idea

## ✗ Without RMDP

Different home/work → Different graph

- Alice (Zone 50 → Zone 73): Build graph #1
- Bob (Zone 12 → Zone 88): Build graph #2
- Carol (Zone 99 → Zone 25): Build graph #3
- ...

= 20,736 separate graphs!

## ✓ With RMDP

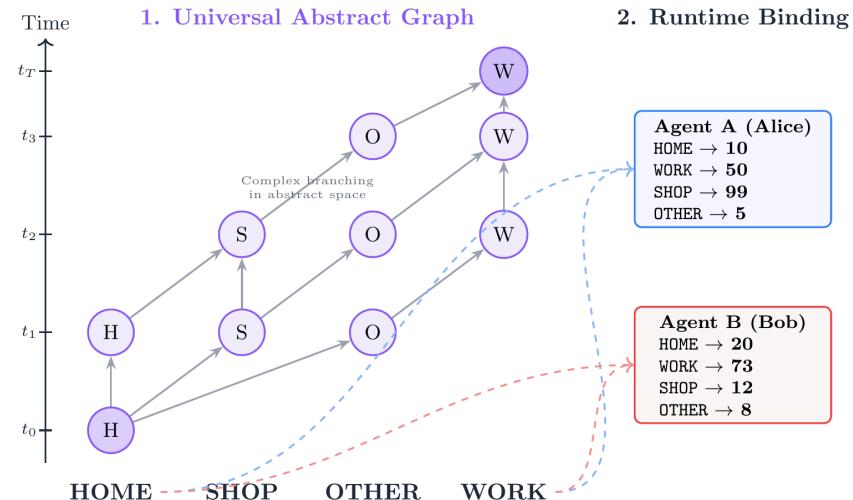
Use **abstract roles** instead of zones

- Graph uses: HOME, WORK, SHOP
- Alice binds: HOME=50, WORK=73
- Bob binds: HOME=12, WORK=88
- Same graph, different bindings!

= 1 universal graph!

Think of it like a map template: "Go from HOME to WORK" is the same instruction for everyone. Only the *addresses* differ.

# RMDP: How Zone Binding Works



#### Same Graph, Different Reality:

When the graph says "Go to WORK", Alice goes to Zone 50, Bob goes to Zone 73. This allows **millions of agents** to share a single graph structure in GPU memory.

# Why This Works: Same Logic, Different Numbers

## The Logic (Same for Everyone)

- "Leave home in the morning"
- "Travel to work"
- "Maybe stop for shopping on the way back"
- "Return home by evening"

## The Numbers (Different for Everyone)

- Home → Work travel time: 25 min vs 40 min vs 15 min
- Which shops are nearby
- Specific zone preferences

RMDP separates **WHAT** decisions are made (graph structure) from **WHERE** they happen (zone binding)

Result: 316 hours → 47 seconds

# 5. Results

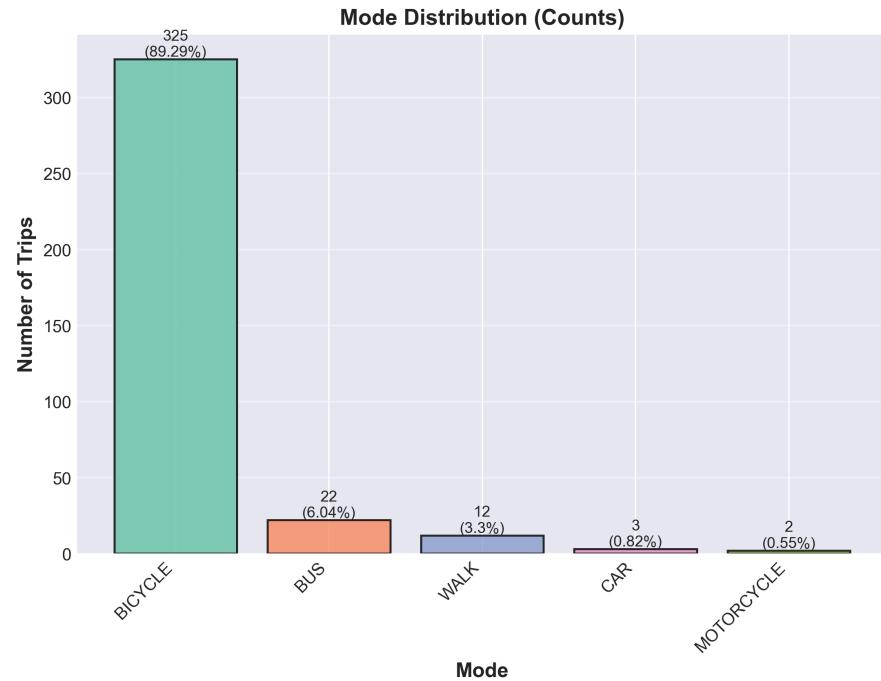
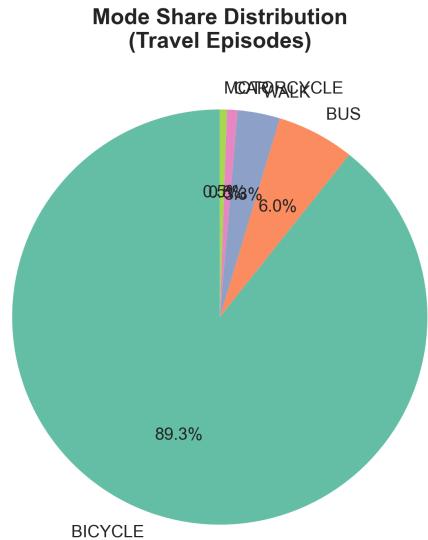
Performance and Validation

# Performance Comparison

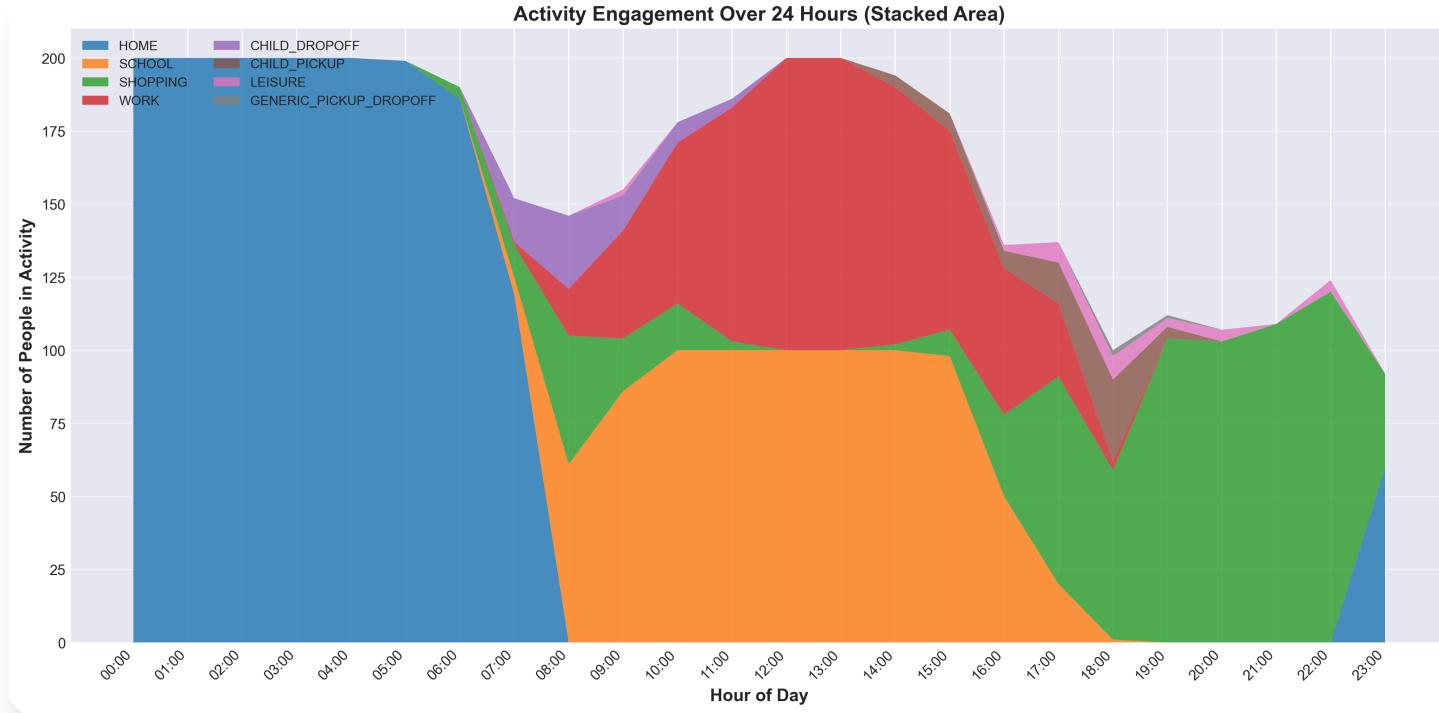
Approach	Time	Speedup
No Pruning (Baseline)	71.1 hours	1×
Dictionary-Based CPU	108s	2,370×
Tensor-Based CPU	65s	3,940×
GPU CUDA	8s	32,000×
+ Universal Graph	7s	24.7× vs per-home
+ RMDP (Final)	8.7s	29,382×

From 71 hours to 9 seconds — while maintaining optimality!

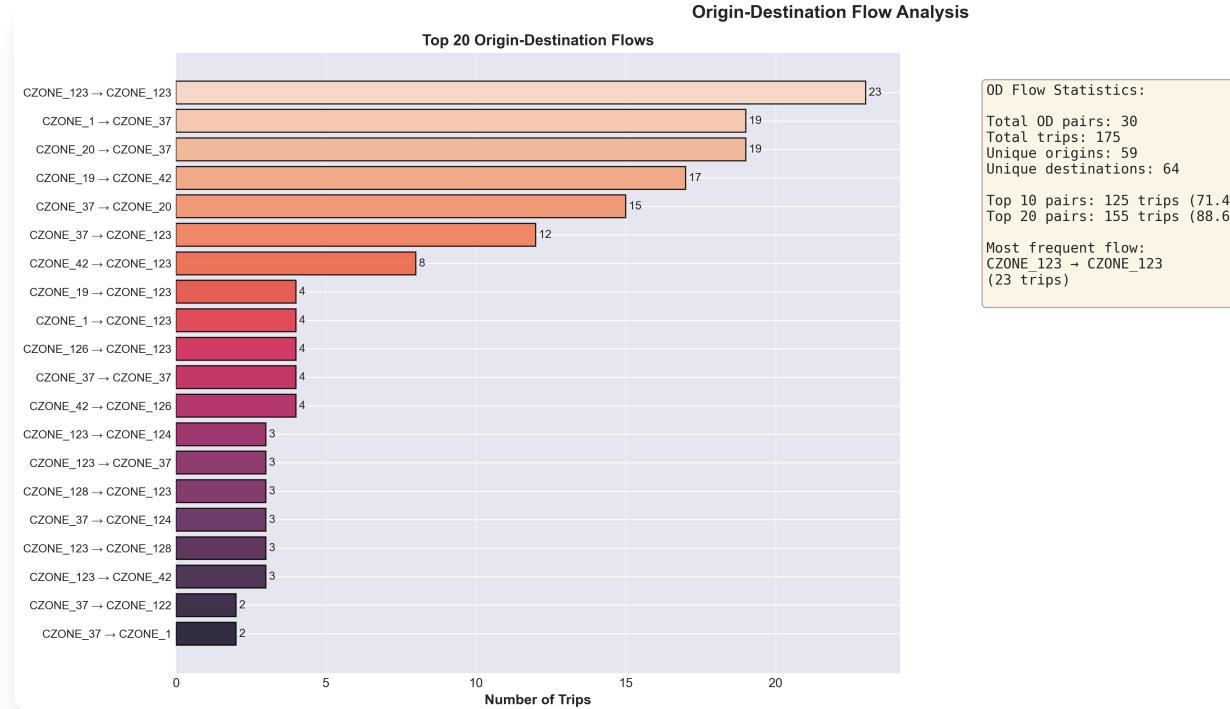
# Simulation Results: Mode Distribution



# Simulation Results: Daily Activity Pattern



# Simulation Results: Travel Flows



# 6. Conclusion

Contributions and Future Work

# Research Contributions

## 1. INTEGRATION

First application of reachability analysis to discrete choice models

## 2. PERFORMANCE

29,382× speedup via RMDP + GPU

## 3. THEORETICAL BRIDGE

Space-Time Prism  $\equiv$  FRT  $\cap$  BRT

## 4. PRACTICAL IMPACT

Enables population-scale simulation on consumer hardware

This work solves the simulation bottleneck, enabling estimation and traffic assignment integration.

# Future Work: The Conditions Framework

**Key Insight:** Reachability analysis generalizes beyond simulation to handle *any constraint* expressible as "what states are reachable?"

## Household Equity

Track travel burden within households

Gender equity under congestion pricing

## Joint Activities

Unified abstraction for care activities

Childcare, eldercare, shared meals

## Work Flexibility

Evaluate policy impacts

4-day week, flextime, remote work

**Same computational foundation** (this work) powers all applications via composable constraint specification.

# Thank You

**Solving the Curse of Dimensionality in DDCM**

Graph-Based Computation | PyTorch & CUDA | CSR Graphs | Level-Sync BFS

29,382× speedup | 71 hours → 9 seconds | Optimal solutions preserved

# Appendix

Technical Details (Backup for Q&A)

# Appendix: The Bellman Equation

The value function can be defined recursively (Bellman 1957, Rust 1987)

$$\bar{V}(x_k) = \log \sum_{a_k \in C(x_k)} \exp [u(x_k, a_k) + EV(x_k, a_k)]$$

## COMPONENTS

- $u(x,a)$ : Immediate utility of action a in state x
- $EV(x,a)$ : Expected future value after taking action a
- **log-sum-exp**: Aggregates over all possible actions

## PROPERTIES

- Gumbel errors → Logit choice probabilities
- Log-sum = Consumer surplus
- Analytically differentiable

# Appendix: MDP State-Action Space

## STATE $S_K$

$t$  Time of day (96 steps)

$l$  Location (1,240 zones)

$p$  Current activity purpose

$\tau$  Duration in activity

$m$  Mode (car/PT/walk/bike)

$h$  Mandatory activity history

## ACTION $A_K$

$\tilde{d}$  Destination (1,240 zones)

$\tilde{m}$  Mode of transport

$\tilde{p}$  Activity purpose

stay = continue current activity.

# Appendix: GPU Parallelization Details

Level-Synchronous BFS: Process all states at each time step in parallel

CSR: COMPRESSED SPARSE ROW

Dense Matrix 8.7 TB ✗

Edge List (COO) 2.6 GB

CSR 1.3 GB ✓

MEMORY BANDWIDTH

CPU (DDR4) 25 GB/s

GPU (GTX 1080 Ti) 484 GB/s

19x bandwidth advantage!

# Appendix: Utility Function Components

Total Utility = Travel + Activity + Timing

$$u(s, a) = u_{\text{travel}} + u_{\text{start}} + u_{\text{act}} + \varepsilon(a)$$

**u<sub>travel</sub>**

Disutility of travel: time, cost, mode constants

**u<sub>start</sub>**

When/where to start: timing preferences, location attractiveness

**u<sub>act</sub>**

Activity duration: marginal utility with diminishing returns

**ε(a)**

Random utility: i.i.d. Gumbel → Logit probabilities

# Appendix: Approximation Methods

## Basis Function BI

$\$V(s, \text{home}) \approx V_{\text{ref}}(s) + \beta \times \Delta \text{travel\_time}$

- Train on 8 representative zones
- Interpolate for remaining 136
- Accuracy:  $R^2 = 0.982$
- Speedup: 9.3x

## Multi-Fidelity BI

Full BI for some zones, interpolate rest

- Full BI for 15 representatives
- IDW interpolation for rest
- Accuracy: 99.6%
- Speedup: 2.5x

# Appendix: Future Research Directions

Gap	Status	Enabled By
Simulation	✓ Solved (29,382x)	This work
Estimation	Ready to implement	Speedup
Traffic Assignment	Possible to integrate	Speedup
Correlation	Research needed	—
Households	Conditions Framework	Reachability
Long-term	Conditions Framework	Reachability

**Future:** The Conditions Framework generalizes reachability to handle household coordination, joint activities, and work flexibility policies.