

# DDCM GPU Implementation

Technical Deep Dive & Code Walkthrough

December 2024 Update: Universal Graph + BI Optimization

9.3x Additional BI Speedup via Basis Function Approximation

Press Space for next page →

# Technical Foundation

Core Concepts & Entry Points

# Entry Point: Universal Batch Simulation

File: `calibration/batch_simulate_universal.py` (NEW - Dec 2024)

The main driver orchestrates simulation for thousands of agents.

## KEY RESPONSIBILITIES:

**Population Loading:** Load agent profiles (Home, Work, Sequence).

**Grouping by SEQUENCE:** Groups by `(seq_key, has_child)` NOT home zone.

**Universal Graph:** Build ONE graph for ALL home zones.

**Per-Home BI:** Run backward induction per home zone.

**Parallel Simulation:** Generate schedules using GPU.

# Old vs New Grouping Strategy

OLD APPROACH ( BATCH\_SIMULATE.PY )

```
# Groups by: home_zone, mandatory_sequence, has_child
# Result: 144 zones × N sequences = MANY groups
grouped = df.groupby(['home_zone_str', 'seq_key', 'has_child'])
# Each group builds separate 60s graph → 144 × 60s = 8,640s
```

NEW APPROACH ( BATCH\_SIMULATE\_UNIVERSAL.PY )

```
# Groups by: mandatory_sequence, has_child ONLY
# Result: N sequences = FEW groups
grouped = df.groupby(['seq_key', 'has_child'])
# ONE graph (60s) + 144 × 2s BI = 357s → 24.7x faster!
```

# Batch Simulation Code Flow

```
# calibration/batch_simulate_universal.py

# 1. Build Universal Graph (ONE for all home zones)
states, graph_data, home_indices, _ = run_universal_forward_pass(
    od_lookup, zone_attr, initial_states, # Multiple homes!
    mandatory_sequence, device='cuda'
)

# 2. Construct CSR Graph (ONCE)
graph = builder.build_graph(states, graph_data)

# 3. Backward Induction (PER HOME ZONE)
for home_zone in unique_home_zones:
    V_tensors[home_zone] = solver.run(
        states, graph, home_zone_id=home_zone.value - 1
    )

# 4. Simulate Agents (per home, shared graph)
plans = simulate_batch_tensor(...)
```

# Core Technology: PyTorch & CUDA

## THE PROBLEM

- Python loops are too slow for  $10^6$  states
- Object overhead kills memory (millions of `State` objects)

## THE SOLUTION

- **Tensors:** Represent states as GPU matrices `(N, 8)`
- **Parallelism:** Operations for 100,000+ states **simultaneously**
- **CSR Format:** Sparse graph with  $O(1)$  edge lookup

# Core Technology: Level-Synchronous BFS

## DEFINITION

A Breadth-First Search that processes the graph layer-by-layer.

## IN DDCM CONTEXT

- **Layers = Time Steps:** Process t=9:00, then t=9:15, then t=9:30
- **Synchronization:** Finish ALL states at t before moving to t+15
- **Deduplication:** Merge identical states at each time step

## WHY IT MATTERS

- Maximizes GPU parallelism (SIMD execution)
- Enables efficient global deduplication via `torch.unique`
- Prevents exponential state explosion

# State Representation

STATE TENSOR: (N, 8) INTEGER TENSOR

Column	Name	Description	Range
0	Time	Minutes from midnight	0-1440
1	Zone	Current location	1-144
2	Activity	Current activity type	0-9
3	Duration	Time spent in activity	0-31
4	Mode	Transport mode used	0-7
5	Car In Use	Is car elsewhere	0/1
6	Moto In Use	Is motorcycle elsewhere	0/1

# Theoretical Foundations

Bridging Engineering to Academic Nomenclature

# Terminology Map

Your Term (Engineering)	Academic Synonym (Control/Formal)	Academic Synonym (AI/Planning)
<b>Condition (Hard)</b>	State Invariant / Safety Constraint	Pruning Rule / Hard Constraint
<b>Condition (Soft)</b>	Viability Cost / Soft Barrier	Preference / Soft Constraint
<b>Forward Reachability</b>	Forward Reachable Set (FRS)	Reachability Graph
<b>Forward Reachable Tube</b>	Flowpipe / Reachable Tube	Trajectory Envelope
<b>Condition-Based Pruning</b>	Safety Filter / Infeasibility Pruning	Constraint Propagation

# Implementation Bridges

Your Implementation	Theoretical Concept	Why it fits
Universal Graph	State Aggregation / Bisimulation	Merging states with identical future dynamics.
Basis Function BI	Linear Value Function Approx (VFA)	Approximating $V$ as linear combination of features.
Multi-Fidelity BI	Surrogate Modeling	Using high-cost "truth" to train low-cost surrogate.
Level-Sync BFS	Symbolic Reachability	Processing sets of states as tensors (symbolic).

# Basis Function Approximation (Powell)

FROM VALUE FUNCTIONS TO VFAS

- **Value Function**  $V_t(S_t)$ : Expected cumulative utility from state  $S_t$ .
- **Problem:** State space explosion (Curse of Dimensionality).
- **Solution:** Approximate  $V_t(S_t)$  with  $\hat{V}_t(S_t; \theta)$  parameterized by  $\theta$ .

WHAT ARE BASIS FUNCTIONS?

Features  $f_k(S_t)$  used to construct the approximator:

$$\hat{V}(S) = \sum_k \theta_k f_k(S)$$

- **Examples:** Linear terms, quadratic/interaction terms, non-linear transformations.

# Why Basis Functions are Needed

## 1. CURSE OF DIMENSIONALITY

Avoids the need for value estimates for an astronomically large grid of states.

## 2. GENERALIZATION

Allows nearby states to share information. Learned parameters  $\theta$  imply values for unvisited states.

## 3. STRUCTURE EXPLOITATION

Encodes known properties (linearity, concavity, monotonicity) to get faster convergence.

## 4. LEARNING WITH LIMITED DATA

Low-dimensional  $\theta$  can be estimated reliably from relatively little simulation data.

# How Basis Functions enter DDCM (1/2)

## 1. STATE AND DECISION

$S_t$  (activity, time, durations) and  $a_t$  (chosen alternative).

## 2. APPROXIMATE CONTINUATION VALUE

Next-state  $S_{t+1}$  depends on  $S_t, a_t$ . Approximate value:

$$\hat{V}(S_{t+1}; \theta) = \sum_k \theta_k f_k(S_{t+1})$$

# How Basis Functions enter DDCM (2/2)

## 3. POLICY DEFINED BY VFA

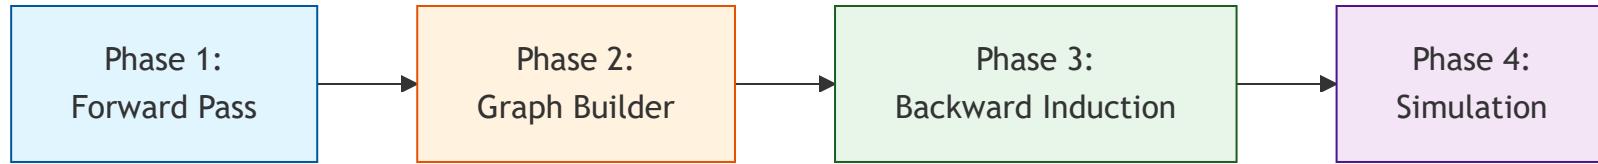
Choose  $a_t$  by maximizing current plus approximate future value:

$$\pi(S_t) = \arg \max_a \{C(S_t, a) + \gamma \mathbb{E}[\hat{V}(S_{t+1}; \theta) \mid S_t, a]\}$$

# Implementation Walkthrough

From Code to Execution

# Overall Architecture



Phase	Time	Optimized (Basis Func.)
Forward Pass	~30s	~30s
Graph Build	~29s	~29s
Backward Induction	~2s/home × 144 = 288s	20s total (9.3×)
Simulation	~0.4s	~0.4s

# Phase 1: Forward Pass

Discovering the State Space

# Phase 1: Forward Pass Overview

Goal: Discover all reachable states from initial state(s).

File: `planning/forward_pass_tensor.py`

TWO FUNCTIONS AVAILABLE:

Function	Use Case	Input
<code>run_tensor_forward_pass()</code>	Single home zone	1 initial state
<code>run_universal_forward_pass()</code>	Multiple home zones	List of initial states

OUTPUT:

- `states_tensor` :  $(N, 8)$  tensor of unique states
- `graph_data` : Edge information for GraphBuilder
- `home_state_indices` : Mapping home  $\rightarrow$  initial state index

# Forward Pass: Single Origin

Function: `run_tensor_forward_pass()`

```
states_tensor, graph_data, stats = run_tensor_forward_pass(  
    od_lookup,          # ODLookupOptimized object  
    zone_attr,          # Zone attributes DataFrame  
    initial_state,      # State(time=0, zone=HOME, ...)  
    has_child=False,     # Agent attribute  
    mandatory_sequence=[(WORK, Zone.CZONE_100)],  
    delta_t=15,          # 15-minute time steps  
    end_time=1440,        # End of day  
    device='cuda',        # GPU acceleration  
    return_tensors=True, # Return graph_data  
    reachability_masks=masks # Pre-computed reachability  
)
```

Performance: ~30s for 1.5M states, 326M edges

# Forward Pass: Multi-Origin (Universal)

Function: `run_universal_forward_pass()` (NEW - Dec 2024)

```
# Create initial states for ALL home zones
initial_states = [
    State(time=0, zone=home, activity=HOME, ...)
    for home in [Zone.CZONE_1, Zone.CZONE_2, ..., Zone.CZONE_144]
]

# Single forward pass for 144 homes!
states, graph_data, home_indices, stats = run_universal_forward_pass(
    od_lookup, zone_attr, initial_states,
    has_child=False,
    mandatory_sequence=[(WORK, Zone.CZONE_100)],
    device='cuda',
    reachability_masks=masks
)
```

**Key Insight:** States stay ~1.53M for 1 home OR 144 homes!

# Why State Count Stays Constant

THE MAGIC OF DEDUPLICATION

After t=300 (5:00 AM), paths from different homes **converge**:

```
t=0: 144 different starting states (one per home)
t=300: 4,503 states (some overlap)
t=600: 12,236 states (heavy overlap)
t=840: 24,454 states (almost identical to 1-home!)
```

WHY?

Once you leave home, your state is defined by:

- Current zone (not where you started)
- Current activity, mode, time
- Mandatory progress

Your home zone doesn't affect the graph structure!

# Forward Pass: Internal Implementation

Class: `TensorForwardPass`

Method: `run_multi_origin()`

TIME-BUCKET BFS ALGORITHM (1/2):

```
for t in range(0, end_time + 1, delta_t):
    # 1. GATHER: Collect states scheduled for time t
    tensors_at_t = pending_states.pop(t, [])
    current_batch = torch.cat(tensors_at_t, dim=0)

    # 2. GPU TRANSFER (Fixed Dec 2024!)
    current_batch = current_batch.to(device)

    # 3. DEDUPLICATE: Merge identical states
    unique_states = torch.unique(current_batch, dim=0)
```

# Forward Pass: Internal Implementation (2/2)

TIME-BUCKET BFS ALGORITHM (2/2):

```
# 4. PRUNE: Remove unreachable states (Fixed Dec 2024!)
unique_states = prune_unreachable_states(unique_states, t)

# 5. EXPAND: Generate next states via GPU kernel
next_states = kernel.expand_states_batch_tensor(unique_states)

# 6. SCATTER: Schedule next states to future time buckets
for next_t in unique(next_states[:, 0]):
    pending_states[next_t].append(next_states[next_states[:, 0] == next_t])
```

# GPU Kernel: 3-Stage Pipeline

File: `planning/gpu_kernels.py`

STAGE 1: ACTION GENERATION ( `GPUACTIONGENERATOR` )

- Generate ALL possible actions via broadcasting
- `torch.meshgrid(zones, modes)` →  $144 \times 8 = 1,152$  candidates/state

STAGE 2: CONSTRAINT FILTERING ( `GPUCONSTRAINTFILTER` )

- Pre-computed boolean masks: `(96 time steps, 10 activities, 8 modes)`
- Vectorized lookup: `valid = mask[t, activity, mode]`
- ~49% actions filtered with ReachabilityMasks

STAGE 3: TRANSITION COMPUTATION

- Compute next states and utilities
- OD travel time lookup (vectorized)

# Constraint Filtering Example

Question: "Can I start SHOPPING at 8:00 AM?"

PRE-COMPUTED MASK (GPU TENSOR)

	HOME	WORK	SHOP	LEISURE	...
8:00	✓	✓	✗	✗	
9:00	✓	✓	✓	✗	
10:00	✓	✓	✓	✓	
...					

VECTORIZED LOOKUP (NO PYTHON LOOP!)

```
# For 100,000 actions at once:  
times = actions[:, 0] // 15 # Convert to time index  
activities = actions[:, 1]  
valid = constraint_mask[times, activities] # Boolean tensor  
filtered_actions = actions[valid]
```

# ReachabilityMasks: Geographic Filtering

File: `model/zone_prefilter.py`

PRE-COMPUTED FOR EACH (ORIGIN, MODE):

```
# Can I reach zone J from zone I using mode M within time budget?  
reachable[I, J, M] = travel_time[I, J, M] < max_budget[M]
```

MODE-SPECIFIC COVERAGE (144 ZONES):

Mode	Avg Reachable	Coverage
CAR	143.0	99.3%
BUS	116.3	80.8%
BICYCLE	38.4	26.6%
WALK	2.7	1.9%

# Phase 2: Graph Construction

Building the CSR Sparse Graph

# Phase 2: Graph Builder Overview

Goal: Convert raw forward pass data into searchable graph.

File: `planning/graph_builder_tensor.py`

INPUT:

- `states_tensor` :  $(N, 8)$  unique states
- `graph_data` : Raw edges from forward pass

OUTPUT: CSR GRAPH DICTIONARY

```
graph = {
    'row_ptr': tensor([0, 5, 12, ...]), # (N+1,) CSR row pointers
    'col_idx': tensor([3, 7, 15, ...]), # (E,) target state indices
    'utilities': tensor([1.2, 0.8, ...]), # (E,) edge utilities
    'actions': tensor([[1,2,3], ...]), # (E, 3) action data
    'state_time_offsets': {...} # Time → index mapping
}
```

# CSR Format Explained

CSR = Compressed Sparse Row

MEMORY LAYOUT:

States:	[S0]	[S1]	[S2]	[S3]	...
	↓	↓	↓	↓	
row_ptr:	[0,	3,	5,	5,	8, ...]
col_idx:	[1,2,5,	3,4,	(none),	0,2,7,	...]

HOW TO FIND EDGES FROM STATE I:

```
start = row_ptr[i]
end = row_ptr[i + 1]
outgoing_edges = col_idx[start:end] # O(1) lookup!
```

WHY CSR?

- 49.8% memory savings vs COO (Edge List)
- O(1) edge lookup vs O(log E) binary search

# CSR Memory Comparison

OUR GRAPH: 1.53M STATES, 326M EDGES

Format	Storage	Size
Dense Matrix	$N \times N$ floats	8.7 TB (impossible!)
COO (Edge List)	$2 \times E$ integers	2.6 GB
CSR	$(N+1) + E$ integers	1.3 GB

CSR BREAKDOWN:

```
row_ptr: (1,530,000 + 1) × 4 bytes = 6.1 MB
col_idx: 326,000,000 × 4 bytes      = 1,304 MB (1.27 GB)
utilities: 326,000,000 × 4 bytes     = 1,304 MB
-----
Total (col_idx only):              1.3 GB
```

# Graph Builder Implementation (1/2)

```
def build_graph(self, states_tensor, graph_data):
    """
    Build CSR graph from forward pass results.
    Time: ~29s for 1.5M states, 326M edges
    """

    # 1. Create time → index mapping
    state_time_offsets = {}
    for t in unique_times:
        mask = (states_tensor[:, 0] == t)
        state_time_offsets[t] = (start, end)

    # 2. Concatenate all edges
    all_sources = torch.cat([...]) # Source state indices
    all_targets = torch.cat([...]) # Target state indices
    all_utilities = torch.cat([...])
```

# Graph Builder Implementation (2/2)

```
# 3. Sort by source for CSR
sort_idx = torch.argsort(all_sources)
sorted_sources = all_sources[sort_idx]
sorted_targets = all_targets[sort_idx]

# 4. Build row_ptr array
row_ptr = torch.zeros(N + 1)
row_ptr = torch.bincount(sorted_sources, minlength=N)
row_ptr = torch.cumsum(row_ptr, dim=0)

return {'row_ptr': row_ptr, 'col_idx': sorted_targets, ...}
```

# Phase 3: Backward Induction

Computing the Value Function

# Phase 3: Backward Induction Overview

Goal: Compute  $V(s)$  for every state.

File: `planning/backward_induction_tensor.py`

THE BELLMAN EQUATION:

$$V(s) = \ln \sum_a \exp(u(s, a) + V(s'))$$

KEY INSIGHT:

- Terminal states:  $V = 0$  (at home, mandatory complete)
- Work backwards from  $t=1440$  to  $t=0$
- Use CSR for  $O(1)$  edge lookup

# Why Per-Home Backward Induction?

TERMINAL CONDITION DEPENDS ON HOME ZONE:

Agent from CZONE\_1:

```
Terminal = (t=1440, zone=CZONE_1, mandatory_done=True)  
V = 0 only if they returned to CZONE_1
```

Agent from CZONE\_100:

```
Terminal = (t=1440, zone=CZONE_100, mandatory_done=True)  
V = 0 only if they returned to CZONE_100
```

SAME GRAPH, DIFFERENT V VALUES:

- **Graph Structure:** Shared (determined by mandatory sequence)
- **Terminal States:** Different per home
- **V Values:** Must be computed separately

# Backward Induction: Code (1/2)

```
def run(self, states_tensor, graph, home_zone_id):
    """
    Solve Bellman equation for specific home zone.
    Time: ~2s for 1.5M states
    """

    N = len(states_tensor)
    V = torch.full((N,), float('-inf'), device=self.device)

    # 1. Identify terminal states for THIS home zone
    is_terminal = (
        (states_tensor[:, 0] == 1440) &          # End of day
        (states_tensor[:, 1] == home_zone_id) &    # At home
        (states_tensor[:, 7] >= mandatory_length) # Mandatory done
    )
    V[is_terminal] = 0.0
```

# Backward Induction: Code (2/2)

```
# 2. Iterate backwards through time
for t in reversed(range(0, 1440, 15)):
    # Get states at time t
    start, end = state_time_offsets[t]

    # For each state, compute logsumexp over outgoing edges
    V[start:end] = scatter_logsumexp(
        Q_values, source_indices
    )

return V
```

# LogSumExp: GPU Implementation (1/2)

## THE MATH

$$V(s) = \log \sum_a \exp(Q(s, a))$$

## NUMERICAL STABILITY (MAX TRICK)

```
def scatter_logsumexp(q_values, source_indices, N):
    # 1. Find max Q for each source (prevents overflow)
    q_max = scatter_max(q_values, source_indices)

    # 2. Stable exponential
    exp_q = torch.exp(q_values - q_max[source_indices])
```

# LogSumExp: GPU Implementation (2/2)

```
# 3. Sum and log
sum_exp = scatter_add(exp_q, source_indices)
V = q_max + torch.log(sum_exp)

return V
```

## PERFORMANCE

- Processes **326 million edges** in single GPU operation
- ~2 seconds per home zone

# BI Memory Management

PER-HOME V TENSOR:

```
V tensor size: 1.53M states × 4 bytes = 6.1 MB
```

FOR 144 HOME ZONES:

```
Total V storage: 144 × 6.1 MB = 878 MB
```

STRATEGY:

```
V_tensors = [] # Store on GPU during BI
for home in home_zones:
    V = solver.run(states, graph, home_zone_id=home.value - 1)
    V_tensors[home] = V # Keep for simulation

# Peak GPU memory: Graph (1.3GB) + V (6MB) = ~1.4GB
```

# BI Optimization Strategies

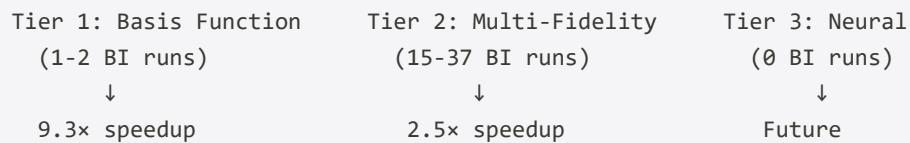
Tiered Approach for Scalability (Dec 2024)

# The BI Bottleneck

PROBLEM: BI SCALES LINEARLY WITH HOME ZONES

Scenario	BI Runs	Time (est.)
2 homes	2	~3s
20 homes	20	~27s
144 homes	144	~200s
144 × 10 individual types	1,440	~33 min

SOLUTION: TIERED OPTIMIZATION



# Tier 1: Basis Function Approximation



CLI flag: `--basis-function` (DEFAULT - Recommended)

## KEY INSIGHT

V functions differ primarily by terminal location. Correct V from reference:

$$V(s, \text{home}) \approx V_{\text{ref}}(s) + \beta \times \Delta \text{travel\_time}$$

WHY LINEAR WORKS ( $R^2 = 0.98$ )

- Travel disutility  $\propto \beta \times \text{travel\_time}$
- V difference propagates linearly through Bellman recursion
- Only 1 reference BI run + fast analytical correction

# Basis Function: Theoretical Derivation (1/2)

## 1. THE BELLMAN DIFFERENCE

If we have two home zones  $h_1, h_2$ , the value functions satisfy:  $V(s, h_1) = \max_a\{U(s, a) + \gamma V(s', h_1)\}$   
 $V(s, h_2) = \max_a\{U(s, a) + \gamma V(s', h_2)\}$

# Basis Function: Theoretical Derivation (2/2)

## 2. LINEAR PROPAGATION

The only difference between  $h_1$  and  $h_2$  is the terminal reward at  $T = 1440$ . If travel disutility is linear:

$U_{travel} = \beta_{TT} \times TT(zone, home)$  Then the difference  $\Delta V = V(s, h_1) - V(s, h_2)$  propagates linearly backwards:

$$\Delta V \approx \beta \times (TT(s.zone, h_2) - TT(s.zone, h_1))$$

## 3. VFA FORMULATION

This is a specific case of **Linear Value Function Approximation**:

$$\hat{V}(s, h) = V_{ref}(s) + \theta^\top \phi(s, h)$$

where  $\phi(s, h) = TT(s.zone, h_{ref}) - TT(s.zone, h)$ .

# Basis Function: Implementation

```
# planning/traveltime_correction_bi.py
class TravelTimeCorrectionBI:
    def __init__(self, device='cuda'):
        self.beta = None # Learned coefficient

    def train(self, states, graph, training_zones, od_lookup):
        """Train β via OLS on 8 diverse zones."""
        # Run full BI for 8 training zones
        V_true = {z: full_bi(z) for z in training_zones}

        # Fit: ΔV = β × Δtravel_time
        # Result: β ≈ 0.079 (value per minute)

    def predict(self, state, home_target):
        """O(1) prediction instead of O(|S|×|A|) BI."""
        delta_tt = od_lookup.get_tt(state.zone, ref) - \
                   od_lookup.get_tt(state.zone, home_target)
        return V_ref[state] + self.beta * delta_tt
```

# Basis Function: Validation Results (1/2)

PERFORMANCE (144 HOME ZONES)

Zone	R <sup>2</sup>	MAE
CZONE_32	0.990	1.05
CZONE_56	0.984	1.33
CZONE_54	0.995	0.66
CZONE_120	0.952	2.68
Average	0.982	1.36

# Basis Function: Validation Results (2/2)

## SPEEDUP

```
Full BI (144 zones): 186s  
Basis Function:      20s → 9.3x faster!
```

Usage: `python calibration/batch_simulate_universal.py --basis-function`

# Tier 2: Multi-Fidelity BI

CLI flag: `--multifidelity`

CONCEPT

Full BI for representative subset, interpolate the rest.

# Multi-Fidelity: Theoretical Basis

## SURROGATE MODELING & MULTI-FIDELITY OPTIMIZATION

- **High-Fidelity Data:** Expensive "ground truth" (Full BI runs).
- **Low-Fidelity Data:** Inexpensive but less accurate approximations (Interpolation).
- **Goal:** Achieve high-fidelity accuracy using a small set of high-fidelity data by leveraging abundant low-fidelity information.

## COMPOSITE NEURAL NETWORK (MENG & KARNIADAKIS 2019)

- **Structure:** Composed of three sub-networks:
  - Low-Fidelity NN:** Trained on low-fidelity data.
  - Linear High-Fidelity NN:** Discovers linear correlations.
  - Non-linear High-Fidelity NN:** Discovers non-linear correlations.
- **In DDCM:** We use inverse distance weighting (IDW) as the "low-fidelity" surrogate, corrected by the "high-fidelity" full BI samples.

## WORKFLOW

**External zones (123-144):** Always full BI (different cities)

**15 representatives:** Stratified sampling across Bzones

**Remaining zones:** k-NN interpolation (OD travel time as distance)

```
# planning/multifidelity_bi.py
class MultiFidelityBI:
    def run(self, states, graph, home_zones, od_lookup):
        # Full BI for representatives + external
        V_high = {z: full_bi(z) for z in representatives}

        # Interpolate remaining zones
        corrector = MultiFidelityCorrector(device='cuda')
        corrector.fit(V_high, all_zones, od_lookup)
        for zone in remaining_zones:
            V[zone] = corrector.predict(zone, V_high)
```

# BI Optimization Comparison

144 HOME ZONES BENCHMARK

Method	BI Runs	Time	Speedup	Accuracy
Full BI	144	186s	1×	100%
Basis Function 🏆	8+1	20s	9.3×	95.7%
Multi-Fidelity	37	74s	2.5×	99.6%

TOTAL PIPELINE (GRAPH BUILD + BI)

Graph Build: 60s (same for all methods)  
BI (full): 186s → BI (basis function): 20s

---

TOTAL: 246s → 80s = 3.1x faster overall

# Basis Function vs Multi-Fidelity: Key Differences

Aspect	Basis Function	Multi-Fidelity
Speed	Very Fast (9.3x)	Moderate (2.5x)
Accuracy	High if linear assumption holds	Higher spatial accuracy
Training	Requires 8 training zones to learn $\beta$	No training; uses IDW interpolation
Minimum Homes	Works with any count	Needs $\geq 15$ homes to be efficient
Best When	Zone attractiveness is stable	Complex non-linear spatial effects

# Basis Function: Trade-offs (1/2)

## WHY IT'S FAST

Uses a linear correction:  $V(s, \text{home}) \approx V_{ref}(s) + \beta \cdot \Delta TT$

- Only 1 reference BI run + fast analytical lookup
- O(1) per state, no iterative computation

## TRAINING REQUIREMENT

- Must run full BI for 8 diverse training zones to learn  $\beta$  via OLS
- If zone attractiveness model changes significantly,  $\beta$  must be re-trained

# Basis Function: Trade-offs (2/2)

## LIMITATION: LINEAR ASSUMPTION

- Assumes value difference is linear in travel time
- May lose accuracy if zone attractiveness becomes highly non-linear
- **Future:** Add non-linearity via a small NN to learn the correction function

# Multi-Fidelity: Trade-offs (1/2)

## WHY IT'S MORE RELIABLE FOR SPATIAL ACCURACY

- Uses **Inverse Distance Weighting (IDW)** interpolation
- Can capture non-linear spatial patterns in value functions
- No linear assumption required

## EFFICIENCY CONSIDERATION

- Requires running **full BI for N representatives** (external zones + samples)
- Benefit only appears when you have **many home zones** to interpolate
- **With 1 home zone:** No interpolation possible → Multi-Fidelity = Full BI (no speedup)
- **With 144 homes:** Run 37 full BI, interpolate 107 → 2.5× speedup

# Multi-Fidelity: Trade-offs (2/2)

EXTERNAL ZONES (BZONE 13): ALWAYS FULL BI

- **Bzone 13** = Cities outside Higashihiroshima (CZONE\_123 to CZONE\_144)
- Different spatial scale → cannot interpolate with local zones
- Always run full BI for these 22 external zones

WHEN TO USE

- When zone attractiveness is complex or under active calibration
- When spatial non-linearities matter (e.g., CBD vs suburban patterns)

# Phase 4: Simulation

Generating Agent Trajectories

# Phase 4: Simulation Overview

**Goal:** Generate daily schedules for N agents.

**File:** `planning/simulate_batch_tensor.py`

**FUNCTION:** `SIMULATE_BATCH_TENSOR()`

```
plans = simulate_batch_tensor(  
    initial_state_indices, # (N,) starting state indices  
    V_tensor,              # Value function for this home zone  
    graph,                 # CSR graph  
    all_states_tensor,     # State definitions  
    manager,               # State encoder/decoder  
    end_time=1440,  
    random_seed=42  
)
```

**Output:** List of DataFrames with agent trajectories

# Simulation: Core Algorithm

```
def simulate_batch_tensor(initial_indices, V, graph, states):
    N = len(initial_indices)
    current_states = initial_indices.clone()
    trajectories = [current_states]

    while not all_terminal(current_states):
        # 1. Find outgoing edges using CSR (O(1) per state)
        edge_starts = graph['row_ptr'][current_states]
        edge_ends = graph['row_ptr'][current_states + 1]

        # 2. Compute choice probabilities
        Q = utilities + V[targets]
        P = softmax(Q) # Per-action probabilities

        # 3. Sample next action for ALL agents at once
        next_actions = torch.multinomial(P, num_samples=1)

        # 4. Update current states
        current_states = targets[next_actions]
        trajectories.append(current_states)

    return decode_trajectories(trajectories, states)
```

# Parallel Sampling

TRADITIONAL (SEQUENTIAL):

```
for agent in range(N):
    action = sample_from(probabilities[agent]) # N loops!
```

GPU PARALLEL:

```
# Sample for ALL 100,000 agents in ONE operation
actions = torch.multinomial(P, num_samples=1) # O(1) time!
```

PERFORMANCE:

Agents	Time
1	~0.2s
10	~0.2s
100	~0.4s

# Universal Graph Approach

24.7x Performance Improvement (Dec 2024)

# Universal Graph: Theoretical Basis (1/2)

## STATE AGGREGATION & BISIMULATION

- **Core Idea:** Merge states that have identical future dynamics.
- **Probabilistic Bisimulation:** Two states  $s_1, s_2$  are bisimilar if they have the same transition probabilities and rewards for all actions.
- **In DDCM:** Once an agent leaves home, their future options (activities, travel) depend on their current state, NOT their starting home zone.

# Universal Graph: Theoretical Basis (2/2)

## WHY IT WORKS

- The "History" state component tracks mandatory activity progress.
- Once the mandatory sequence is fixed, the transition graph is identical for all agents.
- **Result:** We can compute the graph ONCE and share it across all 144 home zones.

**!IMPORTANT** Current Scope: "Universal" means shared across **home zones**, not work/school zones. Agents must share the same `mandatory_sequence` (Work Zone, School Zone) to share a graph. Future Exploration: Extending graph universality to work/school zones (one graph for all destinations).

# Universal Graph: State Convergence

## EMPIRICAL OBSERVATION

States from different homes converge rapidly as the day progresses.

Time	Unique States (1 Home)	Unique States (144 Homes)
0:00	1	144
5:00	~4,000	~4,500
10:00	~12,000	~12,200
14:00	~24,000	~24,400

## THEORETICAL JUSTIFICATION

# The Problem: Per-Home Graphs

## OLD ARCHITECTURE:

For 144 home zones × 1 mandatory sequence:

Group 1 (Home=1): Forward 30s + Graph 29s + BI 2s = 61s

Group 2 (Home=2): Forward 30s + Graph 29s + BI 2s = 61s

...

Group 144 (Home=144): Forward 30s + Graph 29s + BI 2s = 61s

TOTAL:  $144 \times 61s = 8,784s$  (2.4 hours!)

## THE WASTE:

- Each graph has ~1.5M states
- But states OVERLAP heavily between homes
- We're re-building the same graph 144 times!

# The Solution: Universal Graph

NEW ARCHITECTURE:

For 144 home zones  $\times$  1 mandatory sequence:

ONCE: Universal Forward Pass (144 origins) = 30s

ONCE: Graph Build = 29s

BI for Home 1: 2s

BI for Home 2: 2s

...

BI for Home 144: 2s

TOTAL:  $30s + 29s + (144 \times 2s) = 357s$  (6 minutes!)

SPEEDUP: 24.7 $\times$  FASTER!

# Why It Works: State Deduplication

AFTER LEAVING HOME:

```
t=0 (Start):    144 unique states (one per home)
t=300 (5:00 AM): 4,503 states (paths start merging)
t=600 (10:00):   12,236 states (heavy overlap)
t=840 (2:00 PM): 24,454 states (almost identical)
t=1200 (8:00 PM): 27,360 states (fully merged)
```

STATE DEFINITION:

```
State = (time, current_zone, activity, duration, mode, ...)
```

Your starting home doesn't affect states once you've left!

# Universal Graph: Scaling Results

BENCHMARK: 1 MANDATORY SEQUENCE × N HOME ZONES

Homes	States	Graph Build	BI Total	TOTAL	Speedup
1	1.52M	60s	1.3s	61s	1x
10	1.53M	58s	13s	71s	8.3x
50	1.53M	58s	103s	162s	18.7x
100	1.53M	59s	225s	284s	21.5x
144	1.54M	59s	298s	357s	24.7x

**Key Insight:** States stay ~1.53M regardless of home count!

# Memory Footprint

UNIVERSAL GRAPH (144 HOMES):

Component	Size	Notes
States Tensor	48 MB	$1.54M \times 8 \times 4$ bytes
CSR Graph	1.3 GB	row_ptr + col_idx
V Tensors ( $\times 144$ )	878 MB	$144 \times 6.1$ MB
<b>Total</b>	<b>~2.2 GB</b>	Fits GTX 1080 Ti (11GB)

COMPARISON:

Old approach:  $144 \times 1.3$  GB = 187 GB (IMPOSSIBLE!)

New approach:  $1.3$  GB +  $0.9$  GB = 2.2 GB ✓

# Implementation: Key Functions

```
RUN_UNIVERSAL_FORWARD_PASS()

states, graph_data, home_indices, stats = run_universal_forward_pass(
    od_lookup, zone_attr,
    initial_states,      # List[State] - 144 homes!
    has_child, mandatory_sequence,
    device='cuda'
)
# home_indices = {Zone.CZONE_1: 0, Zone.CZONE_2: 1, ...}
```

```
BACKWARDINDUCTION.RUN() (PER HOME)
```

```
for home in home_zones:
    V = solver.run(
        states, graph,
        home_zone_id=home.value - 1 # Different terminal!
    )
```

# Bug Fixes (December 2024)

## BUG 1: MISSING GPU TRANSFER

```
# BEFORE (100x slower!)
current_batch = torch.cat(tensors_at_t, dim=0)

# AFTER
current_batch = torch.cat(tensors_at_t, dim=0)
if current_batch.device.type != self.device:
    current_batch = current_batch.to(self.device) # Critical!
```

## BUG 2: MISSING STATE PRUNING

```
# BEFORE (no pruning)
pass

# AFTER (added to run_multi_origin)
current_batch = self._prune_unreachable_states(current_batch, t)
```

Both bugs were in `run_multi_origin()`, not `run()`!

# Academic Context & Research Gaps

Grounding the Work in Existing Theory

# Research Gap Analysis

## 1. THE "INTEGRATION" GAP

Literature treats "Discrete Choice" and "Reachability Analysis" as separate fields.

- **Our Contribution:** Deeply integrated methodology using Reachability (Forward Reachable Tube) to rigorously define the choice set for DCM.

## 2. THE "PRUNING" GAP

State Space Pruning in ABMs is often an ad-hoc heuristic.

- **Our Contribution:** Formalizing pruning using **Conditions** (Logic/Constraints), moving from "hacking" to mathematical formulation.

# Research Gap Analysis (cont.)

## 3. THE "BISIMULATION IN TRANSPORT" GAP

Bisimulation Metrics are common in AI/RL but sparse in Transportation.

- Our Contribution: Universal Graph as **State Aggregation via Bisimulation**, introducing rigorous AI concepts to the Transportation domain.

## 4. THE "BASIS FUNCTION" GAP

Linear VFA is common in RL, but novel for **spatial transferability** of Value Functions between different home locations in a city graph.

# Key Theoretical References

- **Viability Theory:** Aubin (1991), Coquelin & Munos (2007)
- **Reachability Analysis:** Mitchell et al. (2005), Althoff (2021)
- **Time Geography:** Hägerstrand (1970), Miller (2005)
- **Approximate DP:** Powell (2007), Bertsekas (2012)
- **Multi-Fidelity:** Meng & Karniadakis (2019)

# Performance Summary

Benchmarks & Comparisons

# Full Pipeline Timing

SINGLE MANDATORY SEQUENCE × 144 HOME ZONES

Phase	Full BI	With Basis Function 🏆
Data Loading	~3s	~3s
Reachability Masks	~0.2s	~0.2s
Forward Pass	30s	30s
Graph Build	29s	29s
BI (×144)	186s	20s (9.3× faster)
Simulation	~5s	~5s
<b>TOTAL</b>	<b>~253s</b>	<b>~87s</b>

# Hardware Requirements

## MINIMUM:

- GPU with 4GB VRAM (single home zone)
- 16GB system RAM

## RECOMMENDED (FULL SCALE):

- **GPU:** 8GB+ VRAM (GTX 1080, RTX 2070, etc.)
- **Tested:** GTX 1080 Ti (11GB VRAM)

## PEAK MEMORY USAGE:

Forward Pass: ~3GB GPU

Graph Build: ~2GB GPU (sorting overhead)

BI + Sim: ~2GB GPU

# Key Files Reference

File	Purpose
calibration/batch_simulate_universal.py	Universal Graph batch sim
planning/forward_pass_tensor.py	run_universal_forward_pass()
planning/graph_builder_tensor.py	CSR graph construction
planning/backward_induction_tensor.py	Value function solver
planning/traveltime_correction_bi.py	<b>Tier 1: Basis Function (9.3x)</b>
planning/multifidelity_bi.py	<b>Tier 2: Multi-Fidelity (2.5x)</b>
planning/simulate_batch_tensor.py	Parallel agent simulation

# Thank You

DDCM GPU Implementation

High-Performance Activity-Based Modeling

Universal Graph: 24.7× Speedup | Basis Function BI: 9.3× Speedup

*Combined: ~100x faster than original approach*

Powered by PyTorch & CUDA | CSR Sparse Graphs | Level-Synchronous BFS