

# Solving the Curse of Dimensionality

Graph-Based Computation for Activity-Based Travel Demand Models

## RESEARCH OBJECTIVE

Solving the curse of dimensionality of DP nature on DDCM activity-based travel demand model with  
graph-based computation approach.

Press Space for next page →

# **Part I: Theoretical Foundation**

Understanding the Problem & Our Approach

# The DDCM Framework

Dynamic Discrete Choice Models model sequential decision-making over time.

THE CORE EQUATION (BELLMAN)

$$V(s) = \ln \sum_a \exp(u(s, a) + V(s'))$$

- $V(s)$ : Value of being in state  $s$
- $u(s, a)$ : Immediate utility of action  $a$
- **LogSumExp**: "Soft Max" from Gumbel-distributed errors

# The Curse of Dimensionality

A state is a multi-dimensional tuple:

Dimension	Size	Example
Location	1,240	zones
Mode	8	car, bus, walk...
Activity	10	work, shop, leisure...
Time	96	15-min steps
History	varies	mandatory progress

THE PROBLEM

# Why It's Computationally Intractable

## MEMORY

Storing the utility array for a realistic city requires  
**Terabytes of RAM**

## TIME

Västberg (2020): 4-10s per agent → Full city: **1,000+ days**

## THE DENSE MATRIX TRAP

Traditional Recursive Logit requires matrix inversion:

$$Z = (I - M)^{-1}b$$

Even if  $M$  is sparse, the inverse is **always dense** →  $O(N^3)$  complexity

# Our Objective

## THE PROBLEM

- ✗ Intractable for large cities
- ✗ Requires massive clusters
- ✗ Slow iteration cycles

## TARGET OUTCOME

- ✓ Orders of magnitude faster
- ✓ Graph-based computation
- ✓ Consumer GPU hardware

# The Key Insight

Reachability & Condition-Based Pruning

# Inspiration: Reachability Analysis

From Control Theory and Robotics (Doshi et al., 2022)

We don't calculate paths for every point. We calculate **Reachable Tubes**:

- **BRT:** States that *can* reach the target
- **FRT:** States reachable from start

"Don't solve for the world. Solve for the Tube."




Fig. 1: Target set and backward reachable set. Several trajectories are shown starting at the same time  $t$  but from different states  $x$  and subject to different input signals  $a(\cdot)$  and  $b(\cdot)$ . Input signal  $a(\cdot)$  is chosen to drive the trajectory away from the target set, while input signal  $b(\cdot)$  is chosen to drive the trajectory toward the target. Figure taken from [37].

# Mapping to DDCM

Apply **Conditions** to prune the state space:

## HARD CONSTRAINTS

- **Time:** Can't teleport
- **Geography:** Won't walk 50km
- **Logic:** Can't drive without car

## THE RESULT

$S_{\text{reachable}} \ll S_{\text{naive}}$

**~99.99%**

State Reduction

# Graph Theory Foundations

The Paradigm Shift

# From Matrices to Graphs

## OLD: MATRIX INVERSION

$$Z = (I - M)^{-1} b$$

- Store full transition matrix
- Inversion is  $O(N^3)$
- Dense output

## NEW: GRAPH TRAVERSAL

$$G = (V, E)$$

- Store only reachable states
- BFS is  $O(V + E)$
- Stays sparse

# Level-Synchronous BFS

Process graph layer by layer where layers = time steps

The "Wave" Algorithm:

**Gather** all states at time  $t$

**Deduplicate** identical states (merge paths)

**Expand** to generate next states

**Scatter** to future time buckets

WHY IT WORKS

- Maximizes GPU parallelism (SIMD)
- Prevents exponential explosion via deduplication
- Natural fit for time-dependent problems

# CSR: Compressed Sparse Row

Memory-efficient graph storage

Format                    1.5M States, 326M Edges

---

Dense Matrix            8.7 TB ✗

---

Edge List (COO)        2.6 GB

---

CSR                      1.3 GB ✓

---

## HOW CSR WORKS


Store only non-zero entries using 3 arrays:

- **Row Pointers:** Where each row starts
- **Column Indices:** Destinations

# **Part II: Implementation**

The 4-Phase Pipeline

# Overall Architecture



Phase	Time	What It Does
Forward Pass	~30s	Discover reachable states
Graph Build	~29s	Construct CSR graph
Backward Induction	~2s/home	Compute value function
Simulation	~0.4s	Generate trajectories

# State Representation

State Tensor: (N, 8) integers on GPU

Col	Name	Range
0	Time	0-1440
1	Zone	1-144
2	Activity	0-9
3	Duration	0-31
4	Mode	0-7
5-6	Vehicle In Use	0/1
7	History	0-15

# Backward Induction

Solving the Bellman Equation

$$V(s) = \ln \sum_a \exp(u(s, a) + V(s'))$$

KEY IMPLEMENTATION: SCATTER LOGSUMEXP

- Aggregate Q-values by source state
- Uses "max trick" for numerical stability
- Processes **326M edges** in one GPU operation

Performance: ~2 seconds per home zone

# Optimization 1: Universal Graph

24.7× Speedup

# The Waste in Per-Home Graphs

## OLD APPROACH

144 homes × (Forward 30s + Graph 29s + BI 2s) = 8,784s

## THE PROBLEM

Each graph has ~1.5M states, but they **heavily overlap!**

# Universal Graph Solution

## KEY INSIGHT (BISIMULATION)

Once you **leave home**, your options depend on **current state**, not origin.

Time	1 Home	144 Homes
0:00	1	144
10:00	12,000	12,200
14:00	24,000	24,400

## NEW APPROACH

```
Forward Pass (144 origins): 30s ← ONCE
Graph Build: 29s ← ONCE
BI per home: 144 × 2s = 288s
TOTAL: 357s → 24.7x faster!
```

# Optimization 2: Basis Function

9.3× BI Speedup

# The Theory

LINEAR VALUE FUNCTION APPROXIMATION

$$\hat{V}(s, h) = V_{ref}(s) + \beta \times \Delta TT$$

where  $\Delta TT$  = travel time difference to home

WHY IT WORKS

- $V$  functions differ mainly by **terminal location**
- Difference propagates **linearly** through Bellman
- Train  $\beta$  on 8 zones → predict all 144

Result:  $R^2 = 0.98$ , **9.3x faster**

# Optimization 3: RMDP

20,736× Reduction

# The Scaling Problem

HOME × WORK × SCHOOL COMBINATIONS

Scenario	Graphs Needed	Time
144 homes × 144 works	20,736	<b>13+ days</b>
Universal Graph	144	~3 hours

We need something better...

# RMDP: Relational MDP

CORE IDEA (BOUTILIER ET AL., 2001)

Use **abstract roles** instead of concrete zones:

ZoneID	Type	Description
HOME	Abstract	Agent's home
WORK	Abstract	Agent's work
ZONE_1..144	Concrete	Fixed zones

The decision structure is identical for all agents!

# RMDP in Action

## ZONEBINDING: RESOLVE AT RUNTIME

State: "I'm at WORK at 5 PM"

Alice (WORK = Zone 50) → resolves to Zone 50

Bob (WORK = Zone 73) → resolves to Zone 73

One graph works for ALL (home, work) combinations

# RMDP Performance

Metric	Before	After	Improvement
Graphs needed	20,736	1	<b>20,736×</b>
Build time	13+ days	~47s	~24,000×
Memory	$O(N \text{ graphs})$	$O(1)$	Massive

# Results Summary

# Full Pipeline Comparison

144 HOME ZONES × 1 MANDATORY SEQUENCE

Phase	Original	+ Universal	+ Basis Func
Forward	144×30s	30s	30s
Graph	144×29s	29s	29s
BI	288s	288s	20s
<b>TOTAL</b>	<b>8,784s</b>	<b>347s</b>	<b>79s</b>
<b>Speedup</b>	<b>1×</b>	<b>25×</b>	<b>111×</b>

# RMDP: The Ultimate Scaling

144 × 144 HOME × WORK COMBINATIONS

Approach	Graphs	Time
Original	20,736	13+ days
RMDP	1	~47s

# Research Contributions

**Integration Gap:** Discrete Choice + Reachability Analysis

**Bisimulation Gap:** Universal Graph via State Aggregation

**Basis Function Gap:** Travel-time VFA for spatial transfer

**RMDP Gap:** Role abstraction in activity-based modeling

# Summary

**24.7×**

Universal Graph

**9.3×**

Basis Function

**20,736×**

RMDP

Combined: ~100× - 20,000× faster

From 13 days → 47 seconds

# Future Direction: Generalized Conditions

Beyond optimization → Generalization

THE VISION: PLUG-AND-PLAY CONSTRAINTS

Make conditions modular so new constraints can be added without rewriting the algorithm

## EXAMPLE EXTENSIONS

- Joint activities (with household)
- Social coordination constraints
- Multi-agent interactions

## IDEAS BEING EXPLORED

- **LLM as Input:** Natural language → constraints
- **Inverse-RL:** Learn constraints from observed data

# Thank You

Solving the Curse of Dimensionality in DDCM

Graph-Based Computation | PyTorch & CUDA | CSR Graphs | Level-Sync BFS